

AMD DEVELOPER HACKATHON: ACT II · TRACK 2

TextSink

Four voices. One truth.

Built end-to-end on Gemma 4 · Fireworks AI

Captioners fail one of two ways.

Accurate but flat.

Frame-bots describe the scene, then miss the joke.

Funny but wrong.

Prompt-bots write a punchline about things that never happened.

The judge scores accuracy **and** tone. TextSink is built to win both.

One clip in. Four voices out.

Official judge clip · all four captions written live by Gemma 4

3



An orange tabby kitten walks forward through the foliage toward the camera in a wooded area.



A fierce predator emerges from the brush, clearly ready to conquer the entire forest one tiny, uncoordinated step at a time.



\$ The latest AI agent successfully navigating its first training environment. It's small, but the feature set is looking promising.



A tiny orange explorer makes a grand entrance, bravely navigating the treacherous jungle of backyard leaves.

How it works

clip → scene facts → four tone contracts → captions

Gemma 4 does every creative step — sees the frames, extracts strict-JSON facts, writes all four styles, judges its own drafts. Gemma-first; a non-Gemma fallback tier exists for harness survival (verified in repo) and never fired in the dress rehearsal.

Same grounding, three products:

results.json	the scored Track 2 contract
TextSink CC	timed closed captions, four voices
The Hecklers	two models arguing about the video

It refuses to caption blind.

Captions may only use extracted facts.
Empty grounding → retry → refuse →
reroute. Empty output blocked by schema.
Fallback chain verified in the repo.

15 clips · 199s · 0 refusals · $\frac{1}{3}$ of budget

A different clip. Same discipline.

FORMAL

A woman with dark hair in a high bun, wearing a beige jacket over an orange top, sits at a white desk in a modern office, typing at a monitor.

SARCASTIC

Just another edge-of-your-seat thriller: a woman typing like her career depends on it.

HUMOROUS-TECH

\$ Spot the race condition: her fingers and the keyboard are both vying for the highest throughput in a single-threaded typing test.

HUMOROUS-NON-TECH

She's typing so fast the keyboard's wondering if it just got a promotion.

All 15 official clips, four voices each, browsable live: banksythequantlab.github.io/textsink

Real closed captions.

Timed .srt tracks that watch the whole video —
one consistent voice per style, beat by beat.

0:00 A tiny orange predator prepares for total
world domination.

0:03 The terrifying advance begins through the
dense shrubbery.

0:06 Truly, a relentless force of nature
approaches.

0:09 The menace is closing in. Send help.

The menace is gaining ground.

Two models. One argument.

Every line generated live, turn by turn — Gemma 4 vs gpt-oss-120b, each reading the other's last remark.

STAN Ocean's attacking the coast! A watery apocalypse, Gus!

GUS Apocalypse? Your kettle's just doing a light rinse.

STAN The deep blue menace is swallowing the very earth!

GUS Only you think foam is world-ending, Stan.

STAN & GUS sarcastic **LINT vs VIBE** tech **DORIS & PEARL** everyday

11-5

Gemma-designed tone mechanisms vs generic style prompts, head-to-head on the fallback model itself, scored by two independent judges. The design is the value.

SEES

vision grounding of every clip

WRITES

all four styles + the CC tracks

ARGUES

plays lead Heckler against gpt-oss

JUDGES

scores its own drafts, ships the winner

TEACHES

its judged winners trained a student that beats its own prompted output 18-4 (neutral judge)

One 26B MoE model, dedicated Fireworks deployment. ***Gemma teaches Gemma.***

Methodology + raw data in the repo: GEMMA_PROVENANCE.md · eval/ABLATIONS.md · ab_results.json



TextSink

Four voices. One truth.

github.com/banksythequantLab/textsink

live gallery: banksythequantlab.github.io/textsink

`docker pull ghcr.io/banksythequantlab/textsink` · MIT · linux/amd64 · Gemma 4 on Fireworks AI