

AMD DEVELOPER HACKATHON: ACT II · TRACK 2

TextSink

Four voices. One truth.

Built end-to-end on Gemma 4 · Fireworks AI

Every entry fails one of two ways.

Accurate but flat.

Frame-bots describe the scene, then miss the joke.

Funny but wrong.

Prompt-bots write a punchline about things that never happened.

The judge scores accuracy **and** tone. TextSink wins both.

One clip in. Four voices out.

Official judge clip · all four captions written live by Gemma 4

3



An orange tabby kitten walks forward through the foliage toward the camera in a wooded area.



A fierce predator emerges from the brush, clearly ready to conquer the entire forest one tiny, uncoordinated step at a time.



\$ The latest AI agent successfully navigating its first training environment. It's small, but the feature set is looking promising.



A tiny orange explorer makes a grand entrance, bravely navigating the treacherous jungle of backyard leaves.

How it works

clip → scene facts → four tone contracts → captions

Gemma 4 does every step — sees the frames, extracts strict-JSON facts, writes all four styles, judges its own drafts.

Same grounding, three products:

results.json	the scored Track 2 contract
TextSink CC	timed closed captions, four voices
The Hecklers	two models arguing about the video

It refuses to caption blind.

Captions may only use extracted facts.

Empty grounding → retry → refuse → reroute. Three model tiers.

An empty caption is structurally impossible.

15 clips · **199s** · **0 refusals** in-container, $\frac{1}{3}$ of budget

Four contracts. Four voices. Real output.

FORMAL

An orange tabby kitten walks forward through the foliage toward the camera in a wooded area.

SARCASTIC

A fierce predator emerges from the brush, ready to conquer the forest one tiny, uncoordinated step at a time.

HUMOROUS-TECH

\$ The latest AI agent successfully navigating its first training environment. Small, but the feature set is looking promising.

HUMOROUS-NON-TECH

A tiny orange explorer makes a grand entrance, bravely navigating the treacherous jungle of backyard leaves.

Same grounded facts under every voice — the typography you see is the product's own caption styling.

Real closed captions.

Timed .srt tracks that watch the whole video —
one consistent voice per style, beat by beat.

0:00 A tiny orange predator prepares for total
world domination.

0:03 The terrifying advance begins through the
dense shrubbery.

0:06 Truly, a relentless force of nature
approaches.

0:09 The menace is closing in. Send help.

The menace is gaining ground.

Two models. One argument.

Every line generated live, turn by turn — Gemma 4 vs gpt-oss-120b, each reading the other's last remark.

STAN Ocean's attacking the coast! A watery apocalypse, Gus!

GUS Apocalypse? Your kettle's just doing a light rinse.

STAN The deep blue menace is swallowing the very earth!

GUS Only you think foam is world-ending, Stan.

STAN & GUS sarcastic **LINT vs VIBE** tech **DORIS & PEARL** everyday

\$5,500

addressable with one project

\$2,500 Track 2 · \$3,000 Best Use of Gemma

SEES

vision grounding of every clip

WRITES

all four styles + the CC tracks

ARGUES

plays lead Heckler against gpt-oss

JUDGES

scores its own drafts, ships the winner

TEACHES

its judged winners trained a student that beats it 18–4

One 26B MoE model, dedicated Fireworks deployment. ***Gemma teaches Gemma.***



TextSink

Four voices. One truth.

github.com/banksythequantLab/textsink

MIT · Docker linux/amd64 · Gemma 4 on Fireworks AI